

PAIR RDDS

There are 2 types of RDDs

Basic RDDs Each element is a single object

Pair RDDs Each element is a
Key/Value pair

Basic RDDs

So far, we have only
worked with Basic RDDs

ie. we treat **each record in**
the RDD as a **single object**

Basic RDDs

All our transformations, actions
act on **each record as a whole**

There are 2 types of RDDs

Basic RDDs Each element is a single object

Pair RDDs Each element is a
Key/Value pair

Pair RDDs

Each element is a
Key/Value pair

Many data processing tasks can be
easily expressed using Key, Value pairs

Ex: Delays by Airline, Sales by City, Word
Counts etc

Pair RDDs

Pair RDDs are special
RDDs where each record
is treated as tuple

Pair RDDs

All the basics RDD transformations
and actions work for Pair RDDs too

Special Transformations and
Actions exist for Pair RDDs

Pair RDDs

Transformations

keys
values
mapValues
groupByKey
reduceByKey
combineByKey

A few
transformations
for pair RDDs

Pair RDDs

Transformations

keys
values

mapValues

groupByKey

reduceByKey

combineByKey

**Return RDDs
with only the
keys or the values**

Pair RDDs

Transformations

mapValues

keys
values

**Takes a function
and applies it on the
values of the key,
value pairs**

groupByKey
reduceByKey
combineByKey

Pair RDDs

Transformations

groupByKey

Groups the values which have the same key into a list/collection

BLR, 3

MUM, 1

BLR, 2



BLR, [3 , 2]

MUM, 1

Pair RDDs

groupByKey

Transformations

cogroup is also like groupByKey

But it can group
values across RDDs

Pair RDDs

Transformations

reduceByKey

This is like reduce on Basic RDDs

It takes a function to
combine 2 values

It combines values
with the same key

Pair RDDs

Transformations

reduceByKey

keys

values

mapValues

groupByKey

combineByKey

Using a Pair RDD of **City, Sales**

you can find **the sum of sales** for each city

Pair RDDs

Transformations

combineByKey

Just as for basic RDDs,
we have **reduce** and
aggregate

For Pair RDDs, we have
reduceByKey and
combineByKey

Pair RDDs

Transformations

combineByKey

keys

values

reduce and aggregate

mapValues

**reduceByKey and
combineByKey**

**Note one important
difference!**

Pair RDDs

Transformations

combineByKey

keys

values

reduce and aggregate

Actions on basic RDDs

mapValues

**reduceByKey and
combineByKey**

**Transformations on
Pair RDDs**

groupByKey

reduceByKey

Pair RDDs

Transformations

One of the most common operations is

to merge 2 Pair RDDs

based on the keys

Pair RDDs

Transformations

Pair RDD1

BLR, 3

MUM, 1


DEL, 2

Pair RDD2

BLR, "B"

MUM, "M"

DEL, "D"



Merge 2 Pair RDDs

BLR, [3, "B"]

MUM, [1, "M"]

DEL, [2, "D"]

Such operations are called joins

Pair RDDs

join

left outer join

right outer join

Transformations

joins

These are similar
to their counter
parts in SQL

Pair RDDs

join

left outer join

right outer join

Transformations

A join will return a new Pair RDD

Values from the input RDDs whose keys match are grouped together

Pair RDDs

Transformations

join

Only keys which exist in both RDDs are returned

Like an inner join
in SQL

left outer join

right outer join

Pair RDDs

Transformations

join

Pair RDD1

Pair RDD2

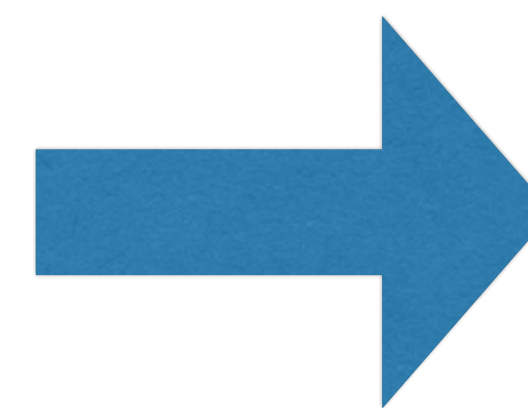
BLR, 3

BLR, "B"

MUM, 1

MUM, "M"

DEL, 2



BLR, [3, "B"]

MUM, [1, "M"]

Pair RDDs

Transformations

join

All keys from the left
RDD are returned

left outer join

right outer join

Pair RDDs

Transformations

left outer join

Pair RDD1

Pair RDD2

BLR, 3

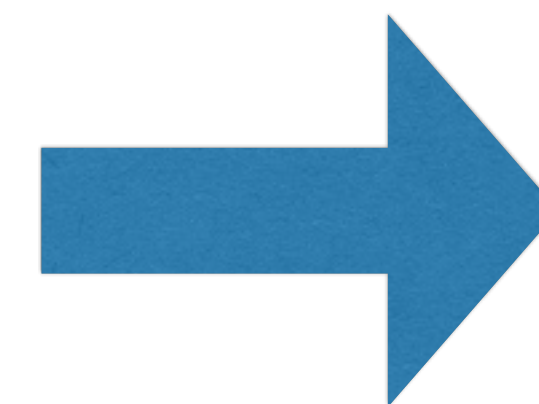
BLR, "B"

MUM, 1

MUM, "M"

DEL, 2

KOL, "D"



BLR, [3, "B"]

MUM, [1, "M"]

DEL, [2, None]

Pair RDDs

Transformations

join

All keys from the right
RDD are returned

left outer join

right outer join

Pair RDDs

Transformations

right outer join

Pair RDD1

Pair RDD2

BLR, 3

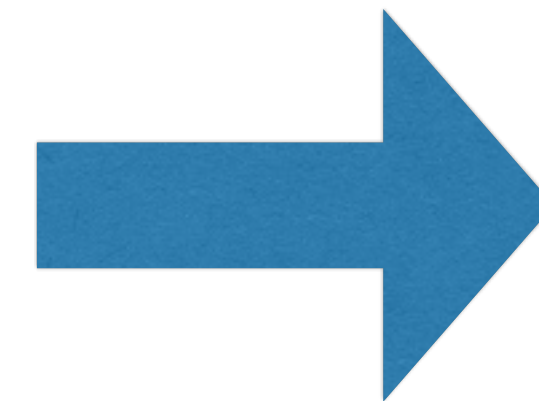
BLR, "B"

MUM, 1

MUM, "M"

DEL, 2

KOL, "D"



BLR, [3, "B"]

MUM, [1, "M"]

KOL, [None, "D"]

Pair RDDs

Actions

countByKey

lookup

collectAsMap

A few special
actions are
available for pair
RDDs

Pair RDDs

Actions

countByKey count the number of values per key

lookup returns all values for a specific key

collectAsMap returns a dict with all the key value pairs

Exploring Airline delays data with PySpark

Part 3

We're back to our Flight
related data from **USDOT**

Let's try doing a few more
things with it

1. Compute the average delay per airport
2. Find the top 10 airports based on delay

1. Compute the average delay per airport

2. Find the top 10 airports based on delay

Average **delay** per airport

For each airport, we need

Sum of all delays

Count of number of flights

Option 1:

Compute these
separately

Option 2:

Compute them
in the same step

Average **delay** per airport

For each airport, we need

Sum of all delays

Count of number of flights

Option 1:

reduceByKey

Option 2:

combineByKey

Recap

We've already created an RDD with Flights data

```
flightsParsed=flights.map(lambda x: x.split(',')).map(parse)
```

Each record is represented as a Flight object

```
Flight(date=datetime.date(2014, 4, 1), airline=u'19805', flightnum=u'1', origin=u'JFK', dest=u'LAX', dep=datetime(8, 54), dep_delay=-6.0, arv=datetime.time(12, 17), arv_delay=2.0, airtime=355.0, distance=2475.0)
```

Recap

```
' , flightnum=u'1' , origin=u'JFK' , dest=u'LAX' , dep=datetime.datetime(2013, 8, 2, 12, 0, 0) ,  
arr_delay=2.0 , airtime=355.0 , distance=2475.0 )
```

**These represent the origin
and destination airport codes**

Average **delay** per airport

Option 1:
reduceByKey

First let's create a **Pair RDD** with origin airport and delay for each flight

```
airportDelays = flightsParsed.map(lambda x: (x.origin,x.dep_delay))
```


Average **delay** per airport

Option 1:
reduceByKey

```
airportDelays = flightsParsed.map(lambda x: (x.origin,x.dep_delay))
```

To create a Pair RDD, just make
sure **each record is a tuple**

Average **delay** per airport

Option 1:
reduceByKey

```
airportDelays = flightsParsed.map(lambda x: (x.origin,x.dep_delay))
```

To see if this worked , you can try to
access the keys and values of the
pair RDD

Average delay per airport

access the keys and values

Option 1:
reduceByKey

```
airportDelays.keys().take(10)
```

```
[u'JFK',  
u'LAX',  
u'JFK',  
u'LAX',  
u'DFW',  
u'OGG',  
u'DFW',  
u'HNL',  
u'JFK',  
u'LAX']
```

keys() and values() are transformations

```
airportDelays.values().take(10)
```

```
[-6.0, 14.0, -6.0, 25.0, -5.0, 126.0, 125.0, 4.0, -7.0, 21.0]
```

Average delay per airport

Option 1:
reduceByKey

```
airportDelays = flightsParsed.map(lambda x: (x.origin,x.dep_delay))
```

On this RDD, we can use
reduceByKey twice, once for the
sum and again for the count

Average **delay** per airport

Option 1:
reduceByKey

```
airportTotalDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

First, let's compute the
SUM

Average **delay** per airport

Option 1:
reduceByKey

```
airportTotalDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

This is a new RDD

reduceByKey is a
transformation

Average **delay** per airport

Option 1:
reduceByKey

```
=airportDelays.reduceByKey(lambda x,y:x+y)
```

Similar to the **reduce** action

reduceByKey takes a function
that combines 2 elements into 1

Average **delay** per airport

Option 1:
reduceByKey

```
airportDelays.reduceByKey(lambda x,y:x+y)
```

The result is a Pair RDD

Keys = Airports
Values = Sum of Delays

Average **delay** per airport

Option 1:
reduceByKey

```
airportDelays.reduceByKey(lambda x,y:x+y)
```

reduceByKey effectively flattens the Pair RDD

Let's see this visually

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	2
JFK	14
PPG	5
PPG	10
JFK	0
PPG	4

P2

JFK	3
JFK	0
LAX	6
LAX	11
JFK	0
PPG	7

This is the delays RDD

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	2
JFK	14
PPG	5
PPG	10
JFK	0
PPG	4

P2

JFK	3
JFK	0
LAX	6
LAX	11
JFK	0
PPG	7

First, the operation
is performed on
each partition

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	2
JFK	14
PPG	5
PPG	10
JFK	0
PPG	4

Values with the
same key are
grouped together

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	2
JFK	14
JFK	0

PPG	5
PPG	10
PPG	4

The given function
is applied **on each**
of these groups

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

p1

JFK	2	<i>x</i>
JFK	14	<i>y</i>
JFK	0	

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

p1

JFK	16
JFK	0

$x + y$


```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

p1

JFK	16	x
JFK	0	y

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	16
-----	----

$x + y$

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	16
-----	----

The same thing is
done for each key

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	16
PPG	5

The same thing is
done for each key

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	16
PPG	5

**We have a set of
Key value pairs
from each node**

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	16
PPG	5

P2

JFK	3
LAX	17
PPG	7

These are **shuffled** so that all the **values with same key** are on a single partition

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

shuffle

P1


JFK	16
PPG	5

JFK	16
JFK	3

P2

JFK	3
LAX	17
PPG	7

PPG	5
PPG	7
LAX	17




```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	16
JFK	3

P2

PPG	5
PPG	7
LAX	17

On each partition,
again the
reduceByKey
operation is applied

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

P1

JFK	19
-----	----

P2

PPG	12
LAX	17

On each partition,
again the
`reduceByKey`
operation is applied

```
alDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

Old RDD

P1

JFK	2
JFK	14
PPG	5
PPG	10
JFK	0
PPG	4

P2

JFK	3
JFK	0
LAX	6
LAX	11
JFK	0
PPG	7

New RDD

P1

JFK	19
-----	----

P2

PPG	12
LAX	17

Average **delay** per airport

Option 1:
reduceByKey

```
airportTotalDelay=airportDelays.reduceByKey(lambda x,y:x+y)
```

We can use the same
idea to compute count

```
airportCount=airportDelays.mapValues(lambda x:1).reduceByKey(lambda x,y:x+y)
```

```
airportCount=airportDelays.mapValues(lambda x:1).reduceB
```

This is our RDD with
(origin airport, delay)

```
rtDelays.mapValues(lambda x:1).reduceByKey(lambda x,y:x+y)
```

(origin airport, delay)

mapValues will leave the
key as is and apply a
function on the value

```
rtDelays.mapValues(lambda x:1).reduceByKey(lambda x,y:x+y)
```

(origin airport, delay) \longrightarrow (origin airport, 1)

All the values are
mapped to the
constant 1


```
rtDelays.mapValues(lambda x:1).reduceByKey(lambda x,y:x+y)
```

(origin airport, 1)

This will sum all these 1s
effectively giving us **count of
flights by airport**

Average delay per airport

Option 1:
reduceByKey

```
airportTotalDelay=airportDelays.reduceByKey(lambda x,y:x+y)  
airportCount=airportDelays.mapValues(lambda x:1).reduceByKey(lambda x,y:x+y)
```

The sum and count are
in 2 separate RDDs

Average delay per airport

Option 1:
reduceByKey

```
airportTotalDelay=airportDelays.reduceByKey(lambda x,y:x+y)  
airportCount=airportDelays.mapValues(lambda x:1).reduceByKey(lambda x,y:x+y)
```

We can merge these
using a join operation

Average **delay** per airport

Option 1:
reduceByKey

```
airportSumCount=airportTotalDelay.join(airportCount)
```

This will merge the 2 RDDs by
matching values with the same key

Average **delay** per airport

Option 1:
reduceByKey


```
airportSumCount=airportTotalDelay.join(airportCount)
```

Sum

JFK	23
LAX	14
PPG	5

Count

JFK	4
LAX	2
PPG	1



sumCount

JFK	(23,4)
LAX	(14,2)
PPG	(5,1)

Average **delay** per airport

Option 1:
reduceByKey

```
airportSumCount=airportTotalDelay.join(airportCount)
```

sumCount

JFK	(23, 4)
LAX	(14, 2)
PPG	(5, 1)

We need to **divide** the
sum by the count

Average **delay** per airport

Option 1:
reduceByKey

```
airportAvgDelay=airportSumCount.mapValues(lambda x : x[0]/float(x[1]))
```

sumCount

JFK	(23, 4)
LAX	(14, 2)
PPG	(5, 1)

Once again, we can use
mapValues to do this

Average delay per airport

Option 1:
reduceByKey

```
airportTotalDelay=airportDelays.reduceByKey(lambda x,y:x+y)  
airportCount=airportDelays.mapValues(lambda x:1).reduceByKey(lambda x,y:x+y)  
airportSumCount=airportTotalDelay.join(airportCount)
```

```
airportAvgDelay=airportSumCount.mapValues(lambda x : x[0]/float(x[1]))
```

With reduceByKey, it took 3 steps to
compute the average delay per airport

Average **delay** per airport

Option 1:
reduceByKey

```
airportTotalDelay=airportDelays.reduceByKey(lambda x,y:x+y)  
airportCount=airportDelays.mapValues(lambda x:1).reduceByKey(lambda x,y:x+y)  
airportSumCount=airportTotalDelay.join(airportCount)
```

1. Compute the sum in 1 RDD
2. Compute the count in another RDD
3. Join the 2 RDDs

Average **delay** per airport

Option 1:
reduceByKey

```
airportTotalDelay=airportDelays.reduceByKey(lambda x,y:x+y)  
airportCount=airportDelays.mapValues(lambda x:1).reduceByKey(lambda x,y:x+y)  
airportSumCount=airportTotalDelay.join(airportCount)
```

With **combineByKey**, we
can do all this in one step

Average **delay** per airport

For each airport, we need

Sum of all delays

Count of number of flights

Option 1:

reduceByKey

Option 2:

combineByKey

Average **delay** per airport

Option 2:
combineByKey

```
airportSumCount2=airportDelays.combineByKey((lambda value:(value,1)),  
                                              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
                                              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

This is a new RDD

combineByKey is a
transformation

Average **delay** per airport

Option 2:
combineByKey

```
airportSumCount2=airportDelays.combineByKey((lambda value:(value,1)),  
                                             (lambda acc, value: (acc[0]+value,acc[1]+1)),  
                                             (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

combineByKey is similar
to **aggregate**

Average **delay** per airport

Option 2:
combineByKey

```
airportSumCount2=airportDelays.combineByKey((lambda value:(value,1)),  
                                              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
                                              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

combineByKey requires
3 functions

Average delay per airport

Option 2: combineByKey

```
airportSumCount2=airportDelays.combineByKey(lambda value:(value,1)),  
                                     (lambda acc, value: (acc[0]+value,acc[1]+1)),  
                                     (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

createCombiner Function

This initializes a value **when a key**
is first seen within a partition

Average delay per airport

Option 2:
combineByKey

```
airportSumCount2=airportDelays.combineByKey((lambda value:(value,1)),  
                                              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
                                              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

merge Function

This specifies how values with the
same key should be combined

within a partition

Average delay per airport

Option 2:
combineByKey

```
airportSumCount2=airportDelays.combineByKey((lambda value:(value,1)),  
                                             (lambda acc, value: (acc[0]+value,acc[1]+1)),  
                                             (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

mergeCombiners Function

This specifies how the results from each partition should be combined

Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey( (lambda value: (value, 1)),  
               (lambda acc, value: (acc[0] + value, acc[1] + 1)),  
               (lambda acc1, acc2: (acc1[0] + acc2[0], acc1[1] + acc2[1])) )
```

createCombiner

merge

mergeCombiners

With combineByKey we have
very granular control over how
the computation should happen

Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey( (lambda value:(value,1)),  
               (lambda acc, value: (acc[0]+value,acc[1]+1)),  
               (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])) )
```

Let's see this visually

Average **delay** per airport

Option 2: **combineByKey**

```
.combineByKey( (lambda value:(value,1)),  
               (lambda acc, value: (acc[0]+value,acc[1]+1)),  
               (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])) )
```

P1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

P2

JFK	3
JFK	0
LAX	6
LAX	11
JFK	0
PPG	7

This the delays RDD

Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey( ( lambda value: (value, 1) ),  
               ( lambda acc, value: (acc[0] + value, acc[1] + 1) ),  
               ( lambda acc1, acc2: (acc1[0] + acc2[0], acc1[1] + acc2[1]) ) )
```

P1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

P2

JFK	3
JFK	0
LAX	6
LAX	11
JFK	0
PPG	7

First, the operation
is performed **on**
each partition

Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey( (lambda value:(value,1)),  
               (lambda acc, value: (acc[0]+value,acc[1]+1)),  
               (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])) )
```

P1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

We start with
the **first record**

Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey( (lambda value:(value,1)),  
               (lambda acc, value: (acc[0]+value,acc[1]+1)),  
               (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])) )
```

p1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

This is first time the
key **"JFK"** is seen

Average delay per airport

Option 2:
combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

P1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

JFK, (2,1)

The **createCombiner** function
is used to initialize a value

Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey( (lambda value:(value,1)),  
               (lambda acc, value: (acc[0]+value,acc[1]+1)),  
               (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])) )
```

p1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

JFK, (2,1)

The next step depends on
whether it's a new key or a
key that's already seen

Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1]))
```

P1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

JFK, (2,1)

The key **"JFK"** has
already been seen

Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1]))
```

p1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

JFK, (2, 1) → acc

value

The **merge** function is used

Average delay per airport

Option 2:
combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

p1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

JFK, (14, 2)

“PPG” is a new key

Average delay per airport

Option 2:
combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

P1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

JFK, (14,2)

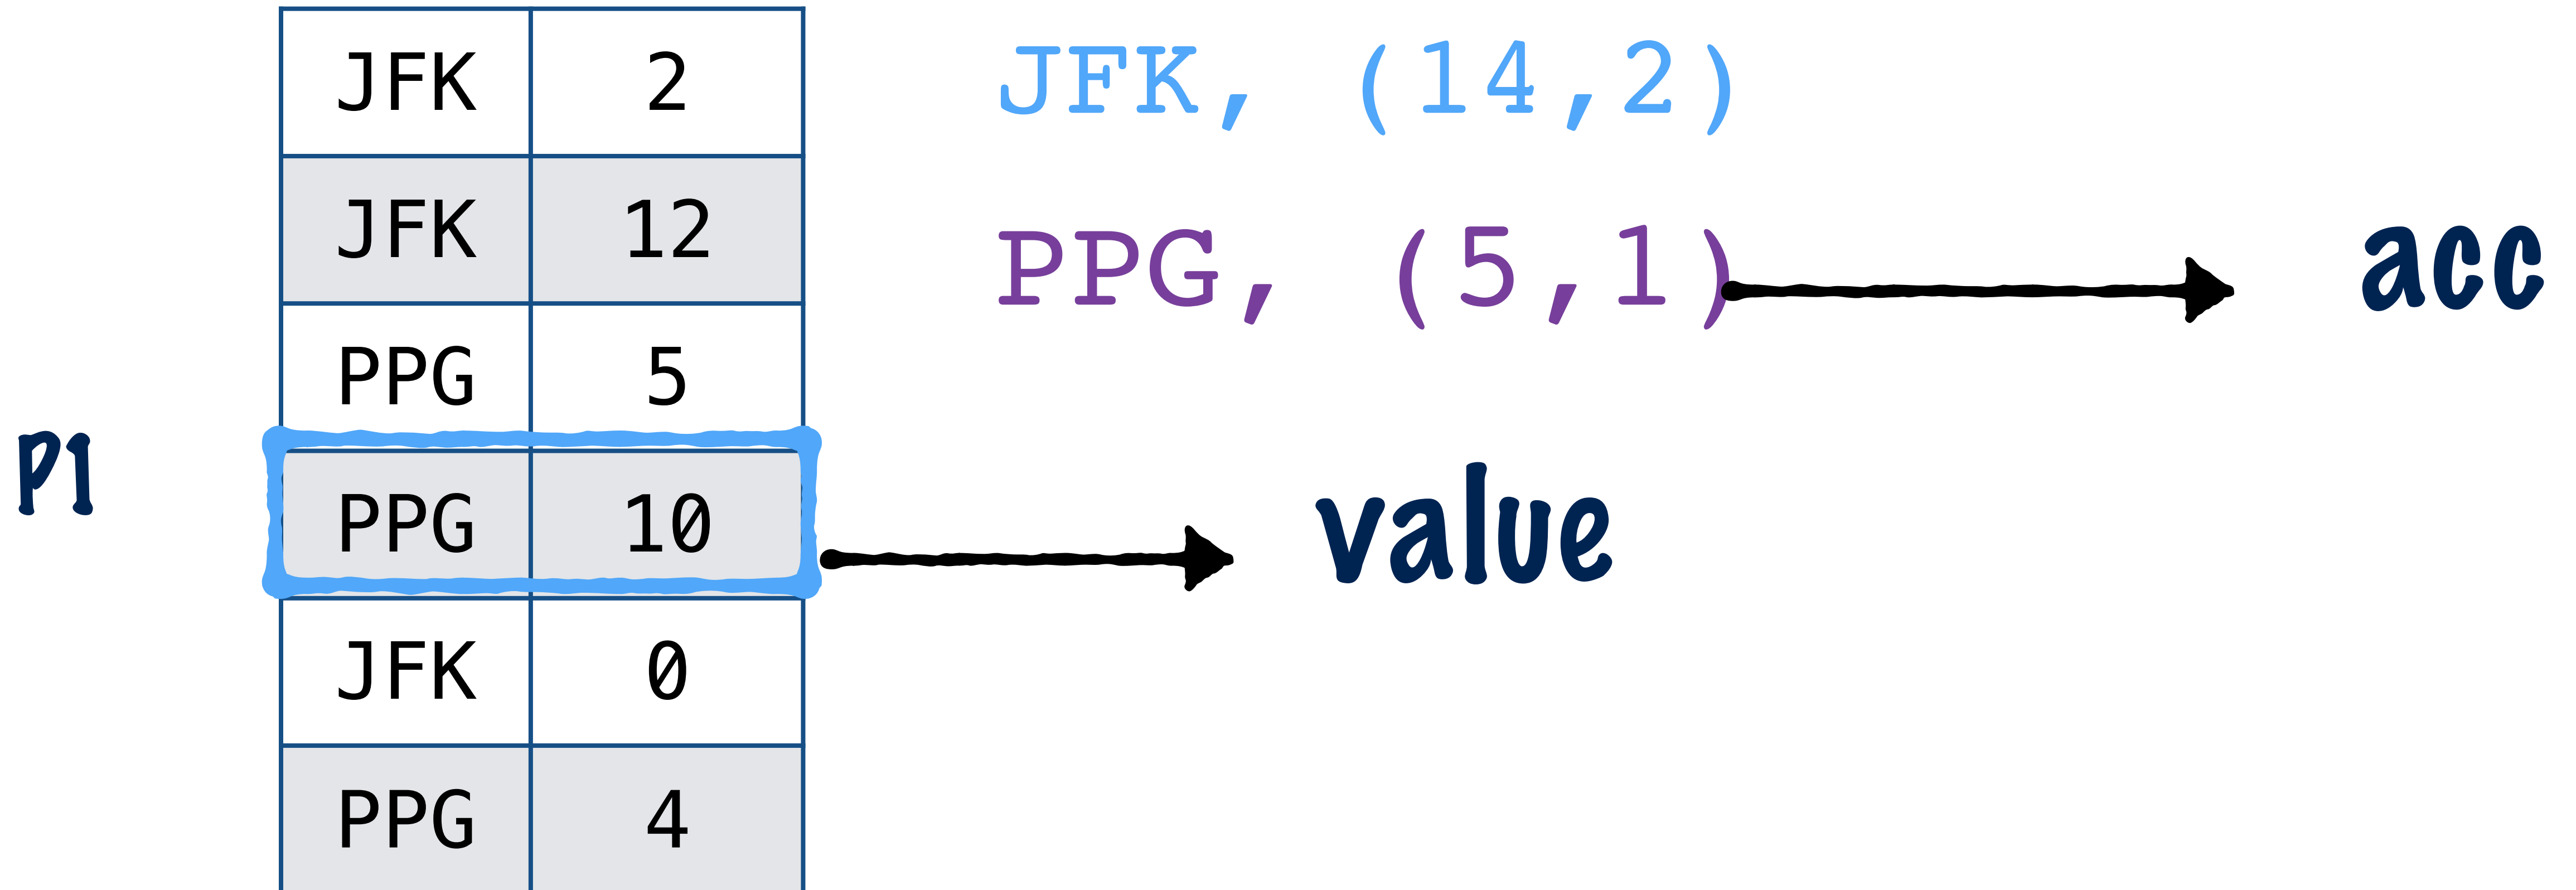
PPG, (5,1)

Continue until all
records are processed

Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1]))
```



Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1]))
```

P1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

JFK, (14, 2) → **acc**

PPG, (15, 2)

→ **value**

Average **delay** per airport

Option 2:
combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1]))
```

p1

JFK	2
JFK	12
PPG	5
PPG	10
JFK	0
PPG	4

JFK, (14, 3)

PPG, (15, 2) → **acc**

→ **value**

Average delay per airport

Option 2: combineByKey

```
.combineByKey( (lambda value:(value,1)),  
               (lambda acc, value: (acc[0]+value,acc[1]+1)),  
               (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])) )
```

p1 JFK, (14,3)
PPG, (19,3)

The same thing
is done on each
Partition

Average delay per airport

Option 2: combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

p1

JFK, (14,3)

PPG, (19,3)

p2

JFK, (3,3)

PPG, (7,1)


LAX, (17,2)

The records are
shuffled until all
records with the same
key are on the same
partition

Average delay per airport

Option 2: combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```



Average delay per airport

Option 2:
combineByKey

```
.combineByKey( (lambda value: (value, 1)),  
               (lambda acc, value: (acc[0]+value, acc[1]+1)),  
               (lambda acc1, acc2: (acc1[0]+acc2[0], acc1[1]+acc2[1])) )
```

P1

JFK, (14, 3)

JFK, (3, 3)

P2

PPG, (19, 3)

PPG, (7, 1)

LAX, (17, 2)

The records in
each partition
are processed

Average delay per airport

Option 2:
combineByKey

```
.combineByKey( (lambda value: (value, 1)),  
               (lambda acc, value: (acc[0]+value, acc[1]+1)),  
               (lambda acc1, acc2: (acc1[0]+acc2[0], acc1[1]+acc2[1])) )
```

p1

JFK, (14, 3)
JFK, (3, 3)


Values with the same
key are combined using
the mergeCombiners
function

Average delay per airport

Option 2:
combineByKey

```
.combineByKey((lambda value:(value,1)),  
              (lambda acc, value: (acc[0]+value,acc[1]+1)),  
              (lambda acc1, acc2: (acc1[0]+acc2[0],acc1[1]+acc2[1])))
```

p1



Average delay per airport

Option 2:
combineByKey

```
.combineByKey( (lambda value: (value, 1)),  
               (lambda acc, value: (acc[0]+value, acc[1]+1)),  
               (lambda acc1, acc2: (acc1[0]+acc2[0], acc1[1]+acc2[1])) )
```

p1 JFK, (17, 6)

Similarly, each
partition is
processed

Average delay per airport

Option 2: combineByKey


```
airportSumCount2=airportDelays.combineByKey((lambda value  
                                              (lambda acc,  
                                              (lambda acc1,
```

Using a single
transformation, we
could compute the sum
and the count

P1 JFK, (17, 6)

P2 PPG, (26, 4)

LAX, (17, 2)



1. Compute the average delay per airport

2. Find the top 10 airports based on delays

The top 10 airports based on delay

```
airportAvgDelay
```

We already have
the average delay
per airport

The top 10 airports based on delay

```
airportAvgDelay
```

It's a pair RDD with

keys=Airports

Values=Avg Delay per Airport

The top 10 airports based on delay

```
airportAvgDelay
```

We want to
sort this RDD

In descending
order of value

The top 10 airports based on delay

```
airportAvgDelay.sortBy(lambda x:-x[1])
```

Descending order of value

sortBy is a transformation

It can be used with both Basic and Paired RDDs

The top 10 airports based on delay

```
airportAvgDelay.sortBy(lambda x: -x[1])
```

This is the avg delay per airport
ie. the value in the key value pair

The top 10 airports based on delay

```
airportAvgDelay.sortBy(lambda x: -x[1])
```

By reversing the sign, we are able
to sort in descending order

The top 10 airports based on delay

```
airportAvgDelay.sortBy(lambda x:-x[1]).take(10)
```

```
[ (u'PPG', 56.25),  
  (u'EGE', 32.0),  
  (u'OTH', 24.533333333333335),  
  (u'LAR', 18.892857142857142),  
  (u'RDD', 18.55294117647059),  
  (u'MTJ', 18.363636363636363),  
  (u'PUB', 17.54),  
  (u'EWR', 16.478549005929544),  
  (u'CIC', 15.931034482758621),  
  (u'RST', 15.6993006993007) ]
```

This will give us the
top 10 airports
and their avg delay

The top 10 airports based on delay

```
[ (u'PPG', 56.25),  
  (u'EGE', 32.0),  
  (u'OTH', 24.533333333333335),  
  (u'LAR', 18.892857142857142),  
  (u'RDD', 18.55294117647059),  
  (u'MTJ', 18.363636363636363),  
  (u'PUB', 17.54),  
  (u'EWR', 16.478549005929544),  
  (u'CIC', 15.931034482758621),  
  (u'RST', 15.6993006993007) ]
```

These are the
airport Codes

Which airports
do these refer to?


The top 10 airports based on delay

```
[ (u'PPG', 56.25),  
  (u'EGE', 32.0),  
  (u'OTH', 24.533333333333335),  
  (u'LAR', 18.892857142857142),  
  (u'RDD', 18.55294117647059),  
  (u'MTJ', 18.363636363636363),  
  (u'PUB', 17.54),  
  (u'EWR', 16.478549005929544),  
  (u'CIC', 15.931034482758621),  
  (u'RST', 15.6993006993007) ]
```

We can use the
airports.csv
file to find out

1. Compute the average delay per airport ✓

2. Find the top 10 airports based on delays ✓



Looking up Airport Descriptions

Recap

There are 3 files

flights.csv

Flight id, airline, airport, departure,
arrival, delay

airlines.csv

airline id, airline name

airports.csv

airport id, airport name

Looking up Airport Descriptions

Let's load and parse
the airports.csv file

```
import csv
from StringIO import StringIO
def split(line):

    reader = csv.reader(StringIO(line))
    return reader.next()

def notHeader(row):
    return "Description" not in row

|
airports=sc.textFile(airportsPath).filter(notHeader).map(split)
```


Looking up Airport Descriptions

```
import csv
from StringIO import StringIO
def split(line):

    reader = csv.reader(StringIO(line))
    return reader.next()

def notHeader(row):
    return "Description" not in row

|
airports=sc.textFile(airportsPath).filter(not
```

A function
to filter out
the header
row

Looking up Airport Descriptions

```
import csv
from StringIO import StringIO
def split(line):
    reader = csv.reader(StringIO(line))
    return reader.next()

def notHeader(row):
    return "Description" not in row

|
airports=sc.textFile(airportsPath).filter(not
```

We take help
from Python's
csv module for a
function to
cleanly parse
the rows

Looking up Airport Descriptions

```
import csv
from StringIO import StringIO
def split(line):
    reader = csv.reader(StringIO(line))
    return reader.next()

def notHeader(row):
    return "Description" not in row
```

We use the functions to
load the data into an RDD
and parse the rows

```
airports=sc.textFile(airportsPath).filter(notHeader).map(split)
```

Looking up Airport Descriptions

```
airports=sc.textFile(airportsPath).filter(notHeader).map(split)
```

airports is also a Pair RDD

A list with 2 elements in each record, this Python treats exactly like a Pair RDD

Looking up Airport Descriptions

```
airports=sc.textFile(airportsPath).filter(notHeader).map(split)
```

airports is also a Pair RDD

Each pair has
(Airport Code,
Airport Description)

Looking up Airport Descriptions

```
airports=sc.textFile(airportsPath).filter(notHeader).map(split)
```

We have 3 options on how to use this RDD

lookup action

map

broadcast

Looking up Airport Descriptions lookup action

```
airports=sc.textFile(airportsPath).filter(notHeader).map(split)
```

Pair RDDs have an
action called **lookup**

Looking up Airport Descriptions lookup action

```
[ (u'PPG', 56.25),  
  (u'EGE', 32.0),  
  (u'OTH', 24.533333333333335),  
  (u'LAR', 18.892857142857142),  
  (u'RDD', 18.55294117647059),  
  (u'MTJ', 18.363636363636363),  
  (u'PUB', 17.54),  
  (u'EWR', 16.478549005929544),  
  (u'CIC', 15.931034482758621),  
  (u'RST', 15.6993006993007) ]
```

We can use
lookup for
each of these
airport Codes

Looking up Airport Descriptions

```
[ (u'PPG', 56.25),  
  (u'EGE', 32.0),  
  (u'OTH', 24.5333333333333335),  
  (u'LAR', 18.892857142857142),  
  (u'RDD', 18.55294117647059),  
  (u'MTJ', 18.363636363636363),  
  (u'PUB', 17.54),  
  (u'EWR', 16.478549005929544),  
  (u'CIC', 15.931034482758621),  
  (u'RST', 15.6993006993007) ]
```

```
airports.lookup('PPG')
```

```
[ 'Pago Pago, TT: Pago Pago International' ]
```

We can use lookup
for each of these
airport Codes

Looking up Airport Descriptions lookup action

```
airports.lookup( 'PPG' )
```

```
[ 'Pago Pago, TT: Pago Pago International' ]
```



**That's an airport on a
tiny little island in the
South Pacific Ocean**

Looking up Airport Descriptions **lookup action**

```
airports.lookup( 'PPG' )
```

```
[ 'Pago Pago, TT: Pago Pago International' ]
```



Did you know that the US has
an island territory there in a
little country called **Samoa**?

Huh!

You learn something
new all the time :)

Looking up Airport Descriptions lookup action

```
airports.lookup( 'PPG' )
```

```
[ 'Pago Pago, TT: Pago Pago International' ]
```

Back to our example,

It's a little tedious to have
to do this for **each airport**

Looking up Airport Descriptions

We have 3 options on how to use this RDD

lookup action ✓

dictionary

broadcast

Looking up Airport Descriptions dictionary

We can **build a map** with all airports from the **RDD** and use that

```
airportLookup=airports.collectAsMap( )
```

collectAsMap is an action

Looking up Airport Descriptions

dictionary

```
airportLookup=airports.collectAsMap( )
```

airportLookup

```
{'SPY': "San Pedro, Cote d'Ivoire: San Pedro Airport",  
'SPZ': 'Springdale, AR: Springdale Municipal',  
'SPP': 'Menongue, Angola: Menongue Airport',  
'SPQ': 'San Pedro, CA: Catalina Air-Sea Terminal Heliport',  
'SPR': 'San Pedro, Belize: San Pedro Airport',  
'SPS': 'Wichita Falls, TX: Sheppard AFB/Wichita Falls Municipal',  
'SPU': 'Split, Croatia: Split Airport',  
'SPW': 'Spencer, IA: Spencer Municipal',  
'SPH': 'Sopu, Papua New Guinea: Sopu Airport',  
'SPI': 'Springfield, IL: Abraham Lincoln Capital',  
'SPK': 'Sapporo, Japan: Chitose AB',  
'SPM': 'Spangdahlem, Germany: Spangdahlem AB',  
'SPN': 'Saipan, TT: Francisco C. Ada Saipan International',  
'SPA': 'Spartanburg, SC: Spartanburg Downtown Memorial',  
'SPB': 'Charlotte Amalie, VI: Charlotte Amalie Harbor Seaplane Base',  
'SPC': 'Santa Cruz de la Palma, Spain: La Palma',  
'SPD': 'Saidpur, Bangladesh: Saidpur Airport',  
'SPE': 'Sepulot, Malaysia: Sepulot Airport',  
'SPF': 'Spearfish, SD: Black Hills Clyde Ice Field',
```

collectAsMap
returns a
dictionary with
all the key value
pairs in the RDD

Looking up Airport Descriptions

dictionary

```
airportLookup=airports.collectAsMap()
```

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

airportAvgDelay is a Pair RDD with

(Airport Code, Avg Delay)

For each record, this function
replaces the Airport Code with
corresponding Description

Looking up Airport Descriptions dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

(Airport Code, Avg Delay)

It uses the **airportLookup**
dict to do so

Looking up Airport Descriptions dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

(Airport Code, Avg Delay)

(Airport Desc, Avg Delay)

airportLookup

```
{'SPY': 'San Pedro, Cote d'Ivoire: San Pedro Airport',  
'SPZ': 'Springdale, AR: Springdale Municipal',  
'SPP': 'Menongue, Angola: Menongue Airport',  
'SPQ': 'San Pedro, CA: Catalina Air-Sea Terminal Heliport',  
'SPR': 'San Pedro, Belize: San Pedro Airport',  
'SPS': 'Wichita Falls, TX: Sheppard AFB/Wichita Falls Municipal',  
'SPU': 'Split, Croatia: Split Airport',  
'SPW': 'Spencer, IA: Spencer Municipal',  
'SPH': 'Sopu, Papua New Guinea: Sopu Airport',  
'SPI': 'Springfield, IL: Abraham Lincoln Capital',  
'SPK': 'Sapporo, Japan: Chitose AB',  
'SPM': 'Spangdahlem, Germany: Spangdahlem AB',  
'SPN': 'Saipan, TT: Francisco C. Ada Saipan International',  
'SPA': 'Spartanburg, SC: Spartanburg Downtown Memorial',  
'SPB': 'Charlotte Amalie, VI: Charlotte Amalie Harbor Seaplane Base',  
'SPC': 'Santa Cruz de la Palma, Spain: La Palma'
```


Looking up Airport Descriptions

dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```



This is a function

Spark distributes this
function to all the nodes with
the partitions of the RDD

Looking up Airport Descriptions

dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

This function uses the
airportLookup variable

It will **carry its own copy**
of this variable to each node

Looking up Airport Descriptions

dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

This is a property of
closure functions

Closure functions are
pretty complicated

Looking up Airport Descriptions

dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

Closure functions are
pretty complicated

But they help Spark take
complex operations defined by
users and apply them on RDDs

Looking up Airport Descriptions dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

But they help Spark take complex operations
defined by users and **apply them on RDDs**

while keeping the user **completely
abstracted** from the complexities
of distributed computing

Looking up Airport Descriptions dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

We mentioned that a copy of airportLookup variable is carried over to each node in the cluster

Looking up Airport Descriptions dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

What if we had to use this
lookup in multiple operations?

Looking up Airport Descriptions dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

The airportLookup variable would be copied over to the nodes every time!

Looking up Airport Descriptions dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

What if we could **cache**
this variable on each of the
nodes, so it can be reused?

Looking up Airport Descriptions dictionary

```
airportAvgDelay.map(lambda x: (airportLookup[x[0]],x[1]))
```

This is exactly why Spark provides

Broadcast Variables

Looking up Airport Descriptions

We have 3 options on how to use this RDD

lookup action ✓

map ✓

broadcast

Looking up Airport Descriptions

broadcast

Spark allows you to define

Broadcast Variables

Looking up Airport Descriptions broadcast

Broadcast Variables

These variables are 'broadcast'
to all the nodes in the cluster

Looking up Airport Descriptions

broadcast

Broadcast Variables

They have 3 characteristics

Immutable

Distributed to all the nodes in the cluster

In-memory

Looking up Airport Descriptions

broadcast

Broadcast Variables

Immutable

Broadcast variables are immutable

They cannot be changed after creation

Looking up Airport Descriptions

broadcast

Broadcast Variables

Distributed to all the nodes in the cluster

in-memory

Looking up Airport Descriptions broadcast

Broadcast Variables

in-memory

Broadcast variables are
cached in memory

So, they should not be too
large

Looking up Airport Descriptions broadcast

```
airportBC=sc.broadcast(airportLookup)
```

This is a broadcast variable

It has a method to lookup the
value for a key

Looking up Airport Descriptions broadcast

```
airportBC=sc.broadcast(airportLookup)
```

```
airportAvgDelay.map(lambda x: (airportBC.value[x[0]],x[1]))
```

This is pretty similar to
how we used the dictionary

Looking up Airport Descriptions **broadcast**

```
lay.map(lambda x: (airportBC.value[x[0]],x[1]))
```

This is pretty similar to how we
used the dictionary

But it's **much more efficient** because
the broadcast variable is cached